

MANIPAL INSTITUTE OF TECHNOLOGY
Manipal – 576 104

DEPARTMENT OF COMPUTER SCIENCE & ENGG.



MANIPAL INSTITUTE OF TECHNOLOGY
MANIPAL
(A constituent unit of MAHE, Manipal)

CERTIFICATE

This is to certify that Ms./Mr.Reg. No.
..... Section: Roll No: has satisfactorily
completed the lab exercises prescribed for Parallel Programming Lab [CSE 3263] of Third Year
B. Tech. Degree at MIT, Manipal, in the academic year 2022-2023.

Date:

Signature
Faculty in Charge

Signature
Head of the Department

CONTENTS

LAB NO.	TITLE	PAGE NO.	REMARKS
	Course Objectives and Outcomes	i	
	Evaluation plan	i	
	Instructions to the Students	ii	
1	Introduction to execution environment of MPI	1	
2	Point to Point Communication in MPI	5	
3	Collective Communication in MPI	9	
4	Collective Communication and Error handling in MPI	12	
5	OpenCL Introduction and programs on Vectors	16	
6	OpenCL programs on Strings, Sorting and to Check the execution time in OpenCL	29	
7	Programs on Arrays in CUDA	35	
8	Programs on Strings in CUDA	44	
9	Programs on Matrix in CUDA	49	
10	Programs on Matrix in CUDA contd.	51	
11	Programs on Matrix in CUDA and image processing applications in CUDA	53	
12	Programs using different CUDA device memory types and synchronization	62	
13	References	66	

Course Objectives

- Learn different APIs used in MPI for point to point, collective communications and error handling
- Learn how to write host and kernel code for device neutral architecture using OpenCL
- Learn how to write host and kernel code in CUDA for nVIDIA GPU card
- To develop the skills of design and implement parallel algorithms using different parallel programming environment

Course Outcomes

At the end of this course, students will be able to

- Write MPI programs using point-to-point and collective communication primitives.
- Solve and test OpenCL programs using GPU architecture
- Develop CUDA programs for parallel applications

Evaluation plan

- Internal Assessment Marks : 60%

➤ Continuous Evaluation : 60%

Continuous evaluation component (for each evaluation):10 marks

The assessment will depend on punctuality, program execution, maintaining the observation note and answering the questions in viva voce.

- End semester assessment of 2 hours duration: 40 %

INSTRUCTIONS TO THE STUDENTS

Pre- Lab Session Instructions

1. Students should carry the Lab Manual Book and the required stationery to every lab session.
2. Be in time and follow the institution dress code.
3. Must Sign in the log register provided.
4. Make sure to occupy the allotted seat and answer the attendance
5. Adhere to the rules and maintain the decorum.
6. Students must come prepared for the lab in advance.

In- Lab Session Instructions

- Follow the instructions on the allotted exercises.
- Show the program and results to the instructors on completion of experiments.
- On receiving approval from the instructor, copy the program and results in the Lab record
- Prescribed textbooks and class notes can be kept ready for reference if required.

General Instructions for the exercises in Lab

- Implement the given exercise individually and not in a group.
- Observation book should be complete with program, proper input output clearly showing the parallel execution in each process. Plagiarism (copying from others) is strictly prohibited and would invite severe penalty in evaluation.
- The exercises for each week are divided under three sets:
 - Solved example
 - Lab exercises - to be completed during lab hours
 - Additional Exercises - to be completed outside the lab or in the lab to enhance the skill
- In case a student misses a lab class, he/ she must ensure that the experiment is completed during the repetition class with the permission of the faculty concerned but credit will be given only to one day's experiment(s).
- Questions for lab tests and examination are not necessarily limited to the questions in the manual, but may involve some variations and / or combinations of the questions.

THE STUDENTS SHOULD NOT

- Bring mobile phones or any other electronic gadgets to the lab.
- Go out of the lab without permission.

Introduction to execution environment of MPI

Objectives:

In this lab, student will be able to

1. Understand the execution environment of MPI programs
2. Learn the various concept of parallel programming
3. Learn and use the Basics API available in MPI

I. Introduction

In order to reduce the execution time work is carried out in parallel. Two types of parallel programming are:

- Explicit parallel programming
- Implicit parallel programming

Explicit parallel programming – These are languages where the user has full control and has to explicitly provide all the details. Compiler effort is minimal.

Implicit parallel programming – These are sequential languages where the compiler has full responsibility for extracting the parallelism in the program.

Parallel Programming Models:

- Message Passing Programming
- Shared Memory Programming

Message Passing Programming:

- In message passing programming, programmers view their programs (Applications) as a collection of co-operating processes with private (local) variables.
- The only way for an application to share data among processors is for programmer to explicitly code commands to move data from one processor to another.

Message Passing Libraries: There are two message passing libraries available. They are:

- PVM – Parallel Virtual Machine
- MPI – Message Passing Interface. It is a set of parallel APIs which can be used with languages such as C and FORTRAN.

Communicators and Groups:

- MPI assumes static processes.
- All the processes are created when the program is loaded.
- No process can be created or terminated in the middle of program execution.
- There is a default process group consisting of all such processes identified by **MPI_COMM_WORLD**.

III. MPI Environment Management Routines:

MPI_Init: Initializes the MPI execution environment. This function must be called in every MPI program, must be called before any other MPI functions and must be called only once in an MPI program.

```
MPI_Init (&argc,&argv);
```

MPI_Comm_size: Returns the total number of MPI processes to the variable size in the specified communicator, such as MPI_COMM_WORLD.

```
MPI_Comm_size(Comm,&size);
```

MPI_Comm_rank: Returns the rank of the calling MPI process within the specified communicator. Each process will be assigned a unique integer rank between 0 and size - 1 within the communicator MPI_COMM_WORLD. This rank is often referred to as a process ID.

```
MPI_Comm_rank Comm,&rank);
```

MPI_Finalize: Terminates the MPI execution environment. This function should be the last MPI routine called in every MPI program. No other MPI routines may be called after it.

```
MPI_Finalize ();
```

Solved Example:

Write a program in MPI to print total number of process and rank of each process.

```
#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[])
{
    int rank,size;

    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    printf("My rank is %d in total %d processes", rank, size);
    MPI_Finalize();
    return 0;
}
```

Steps to execute a MPI program is provided in the form of video which is available in individual systems. However, the basic installation steps are given as follows.

// Follow the following steps to Install, Compile and Run MPI programs in Ubuntu O.S

// To install MPI in Ubuntu, execute the following command in command line

\$sudo apt-get update; sudo apt-get install mpich

// To edit MPI program, use any text editor such as vim or gedit and create a file with .c extension.

//To Compile MPI program, execute the following command in command line

\$mpicc filename.c -o filename.out

//To Run MPI program, execute the following command in command line

\$mpirun -np 4 filename.out

Lab Exercises:

1. Write a simple MPI program to find out pow (x, rank) for all the processes where 'x' is the integer constant and 'rank' is the rank of the process.
2. Write a program in MPI where even ranked process prints "Hello" and odd ranked process prints "World".
3. Write a program in MPI to simulate simple calculator. Perform each operation using different process in parallel.

4. Write a program in MPI to toggle the character of a given string indexed by the rank of the process. Hint: Suppose the string is HeLLO and there are 5 processes, then process 0 toggle 'H' to 'h', process 1 toggle 'e' to 'E' and so on.

Additional Exercises:

1. Write a program in MPI to reverse the digits of the following integer array of size 9 with 9 processes. Initialize the Input array to the following values.
Input array : 18, 523, 301, 1234, 2, 14, 108, 150, 1928
Output array: 81, 325, 103, 4321, 2, 41, 801, 51, 8291
2. Write a MPI program to find the prime numbers between 1 and 100 using two processes.

Point to Point Communication in MPI

Objectives:

In this lab, student will be able to

1. Understand the different APIs used for point to point communication in MPI
2. Learn the different modes available in case of blocking send operation

Point to Point communication in MPI

- MPI point-to-point operations typically involve message passing between two, and only two, different MPI tasks. One task is performing a send operation and the other task is performing a matching receive operation.
- MPI provides both blocking and non-blocking send and receive operations.

Sending message in MPI

- **Blocked Send** sends a message to another processor and waits until the receiver has received it before continuing the process. Also called as **Synchronous send**.
- **Send** sends a message and continues without waiting. Also called as **Asynchronous send**.

There are multiple communication modes used in blocking send operation:

- **Standard mode**
- **Synchronous mode**
- **Buffered mode**

Standard mode

This mode blocks until the message is buffered.

```
MPI_Send(&Msg, Count, Datatype, Destination, Tag, Comm);
```

- First 3 parameters together constitute message buffer. The **Msg** could be any address in sender's address space. The **Count** indicates the number of data elements of a particular type to be sent. The **Datatype** specifies the message type. Some Data types available in MPI are: MPI_INT, MPI_FLOAT, MPI_CHAR, MPI_DOUBLE, MPI_LONG
- Next 3 parameters specify message envelope. The **Destination** specifies the rank of the process to which the message is to be sent.
- **Tag:** The **tag** is an integer used by the programmer to label different types of messages and to restrict message reception.

- **Communicator:** Major problem with tags is that they are specified by users who can make mistakes. **Context** are allocated at run time by the system in response to user request and are used for matching messages. The notions of **context** and **group** are combined in a single object called a communicator (**Comm**).
- The default process group is **MPI_COMM_WORLD**.

Synchronous mode

This mode requires a send to block until the corresponding receive has occurred.

```
MPI_Ssend(&Msg, Count, Datatype, Destination, Tag, Comm);
```

Buffered mode

```
MPI_Bsend(&Msg, Count, Datatype, Destination, Tag, Comm);
```

In this mode a send assumes availability of a certain amount of buffer space, which must be previously specified by the user program through a routine call that allocates a user buffer.

```
MPI-Buffer_attach(buffer, size);
```

This buffer can be released by

```
MPI_Buffer_detach(*buffer, *size);
```

Receiving message in MPI

```
MPI_Recv(&Msg, Count, Datatype, Source, Tag, Comm, &status);
```

- Receive a message and block until the requested data is available in the application buffer in the receiving task.
- The **Msg** could be any address in receiver's address space. The **Count** specifies number of data items. The **Datatype** specifies the message type. The **Source** specifies the rank of the process which has sent the message. The **Tag** and **Comm** should be same as that is used in corresponding send operation. The status is a structure of type status which contains following information: Sender's rank, Sender's tag and number of items received

Finding execution time in MPI

MPI Wtime: Returns an elapsed wall clock time in seconds (double precision) on the calling processor.

- **MPI_Wtime ()**

Solved Example:

Write a MPI program using standard send. The sender process sends a number to the receiver. The second process receives the number and prints it.

```
#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[])
{
    int rank,size,x;
    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Status status;
    if(rank==0)
    {
        Printf("Enter a value in master process:");
        scanf("%d",&x);
        MPI_Send(&x,1,MPI_INT,1,1,MPI_COMM_WORLD);
        fprintf(stdout,"I have sent %d from process 0\n",x);
        fflush(stdout);
    }
    else
    {
        MPI_Recv(&x,1,MPI_INT,0,1,MPI_COMM_WORLD,&status);
        fprintf(stdout,"I have received %d in process 1\n",x);
        fflush(stdout);
    }
    MPI_Finalize();
    return 0;
}
```

Lab Exercises:

- 1) Write a MPI program using synchronous send. The sender process sends a word to the receiver. The second process receives the word, toggles each letter of the word and sends it back to the first process. Both processes use synchronous send operations.
- 2) Write a MPI program where the master process (process 0) sends a number to each of the slaves and the slave processes receive the number and prints it. Use standard send.

- 3) Write a MPI program to read N elements of the array in the root process (process 0) where N is equal to the total number of process. The root process sends one value to each of the slaves. Let even ranked process finds square of the received element and odd ranked process finds cube of received element. Use Buffered send.
- 4) Write a MPI program to read an integer value in the root process. Root process sends this value to Process1, Process1 sends this value to Process2 and so on. Last process sends the value back to root process. When sending the value each process will first increment the received value by one. Write the program using point to point communication routines.

Additional Exercises:

- 1) Write a MPI program to read N elements of an array in the master process. Let N processes including master process check the array values are prime or not.
- 2) Write a MPI program to read value of N in the root process. Using N processes, including root, find out $1! + (1+2) + 3! + (1+2+3+4) + 5! + (1+2+3+4+5+6)$ and print the result in the root process.

Collective Communication in MPI

Objectives:

In this lab, student will be able to

1. Understand the usage of collective communication in MPI
2. Learn how to broadcast messages from root
3. Learn and use the APIs for distributing values from root and gathering the values in the root

Collective Communication routines

When **all processes** in a group participate in a global communication operation, the resulting communication is called a **collective communication**.

MPI_Bcast:

```
MPI_Bcast (Address, Count, Datatype, Root, Comm);
```

The process ranked **Root** sends the same message whose content is identified by the triple (Address, Count, Datatype) to all processes (including itself) in the communicator **Comm**.

MPI_Scatter:

```
MPI_Scatter( SendBuff, Sendcount, SendDatatype, RecvBuff, Recvcount,  
RecvDatatype, Root, Comm);
```

Ensures that the **Root** process sends out personalized messages, which are in rank order in its send buffer, to all the N processes (including itself).

MPI_Gather:

```
MPI_Gather( SendAddress, Sendcount, SendDatatype, RecvAddress, RecvCount,  
RecvDatatype, Root, Comm);
```

The **root** process receives a personalized message from all N processes. These N received messages are concatenated in rank order and stored in the receive buffer of the root process.

Total Exchange:

In routine **MPI_Alltoall()** each process sends a personalized message to every other process including itself. This operation is equivalent to N gathers, each by a different process and in all N^2 messages are exchanged.

Solved Example:

Write a MPI program to read N values of the array in the root process. Distribute these N values among N processes. Every process finds the square of the value it received. Let every process return these value to the root and root process gathers and prints the result. Use collective communication routines.

```
#include "mpi.h"
#include <stdio.h>

int main(int argc, char *argv[])
{
    int rank,size,N,A[10],B[10], c, i;

    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if(rank==0)
    {
        N=size;
        fprintf(stdout,"Enter %d values:\n",N);
        fflush(stdout);
        for(i=0; i<N; i++)
            scanf("%d",&A[i]);
    }
    MPI_Scatter(A,1,MPI_INT,&c,1,MPI_INT,0,MPI_COMM_WORLD);
    fprintf(stdout,"I have received %d in process %d\n",c,rank);
    fflush(stdout);

    c=c*c;
    MPI_Gather(&c,1,MPI_INT,B,1,MPI_INT,0,MPI_COMM_WORLD);

    if(rank==0)
    {
        fprintf(stdout,"The Result gathered in the root \n");
        fflush(stdout);
        for(i=0; i<N; i++)
            fprintf(stdout,"%d \t",B[i]);
        fflush(stdout);
    }

    MPI_Finalize();
    return 0;
}
```

Lab Exercises:

- 1) Write a MPI program to read N values in the root process. Root process sends one value to each process. Every process receives it and finds the factorial of that number and returns it to the root process. Root process gathers the factorial and finds sum of it. Use N number of processes.
- 2) Write a MPI program to read an integer value M and NXM elements into an 1D array in the root process, where N is the number of processes. Root process sends M elements to each process. Each process finds average of M elements it received and sends these average values to root. Root collects all the values and finds the total average. Use collective communication routines.
- 3) Write a MPI program to read a string. Using N processes (string length is evenly divisible by N), find the number of non-vowels in the string. In the root process print number of non-vowels found by each process and print the total number of non-vowels.
- 4) Write a MPI Program to read two strings S1 and S2 of same length in the root process. Using N processes including the root (string length is evenly divisible by N), produce the resultant string as shown below. Display the resultant string in the root process. Use Collective communication routines.

Example:

String S1: string String S2: length Resultant String : slternigtgh

Additional Exercises:

- 1) Write a MPI program to read a value M and NXM number of elements into 1D array in the root, where N is the total number of processes. Find the square of first M numbers, the cube of next M numbers and so on. Print the results in the root.
- 3) Write a MPI program using collective communication functions, to replace all even elements of array A to 1 and replace all odd elements to 0 of size N. Display the resultant array A, count of all even and odd numbers in root process. Assume N is evenly divisible by number of processes.

Example :

Input Array (A): 1 2 3 4 5 6 7 8 9

Resultant Array (A): 0 1 0 1 0 1 0 1 0

Even (Count) = 4

Odd (Count) = 5

Collective Communication and Error Handling in MPI

Objectives:

In this lab, student will be able to

1. Understand the different aggregate functions used in MPI
2. Learn how to write MPI programs using both point to point and collective communication routines
3. Learn and use the APIs for handling errors in MPI

I. Aggregation Functions

MPI provides two forms of aggregation

- **Reduction**
- **Scan**

Reduction:

MPI_Reduce (SendAddress, RecvAddress, Count, Datatype, Op, Root, Comm);

This routine reduces the partial values stored in **SendAddress** of each process into a final result and stores it in **RecvAddress** of the **Root** process. The reduction operator is specified by the **Op** field. Some of the reduction operator available in MPI are: MPI_SUM, MPI_MAX, MPI_MIN, MPI_PROD

Scan:

MPI_Scan (SendAddress, RecvAddress, Count, Datatype, Op, Comm);

This routine combines the partial values into N final results which it stores in the **RecvAddress** of the N processes. Note that root field is absent here. The scan operator is specified by the **Op** field. Some of the scan operator available in MPI are: MPI_SUM, MPI_MAX, MPI_MIN, MPI_PROD

MPI_Barrier(Comm) :This routine synchronizes all processes in the communicator **Comm**. They wait until all N processes execute their respective MPI_Barrier.

Note: All collective communication routines except MPI_Barrier, employ a standard blocking mode of point-to-point communication.

Error Handling in MPI:

- An MPI *communicator* is more than just a group of process that belong to it. Amongst the items that the communicator hides inside is an *error handler*. The error handler is called every time an MPI error is detected within the communicator.
- The predefined default error handler, which is called **MPI_ERRORS_ARE_FATAL**, for a newly created communicator or for MPI_COMM_WORLD is to *abort the whole parallel program* as soon as any MPI error is detected. There is another predefined error handler, which is called **MPI_ERRORS_RETURN**.
- The default error handler can be replaced with this one by calling function **MPI_Errhandler_set**, for example:

```
MPI_Errhandler_set(MPI_COMM_WORLD, MPI_ERRORS_RETURN);
```

- The only **error code** that MPI standard itself defines is **MPI_SUCCESS**, i.e., no error. But the meaning of an error code can be extracted by calling function **MPI_Error_string**. On top of the above MPI standard defines the so called *error classes*. The **error class** for a given error code can be obtained by calling function **MPI_Error_class**.
- Error classes can be converted to comprehensible error messages by calling the same function that does it for error codes, i.e., **MPI_Error_string**. The reason for this is that error classes are implemented as a subset of error codes.

Solved Example:

Write a MPI program using N processes to find $1! + 2! + \dots + N!$. Use collective communication routines.

```
#include <stdio.h>
#include "mpi.h"
int main(int argc, char* argv[])
{
    int rank, size, fact=1, factsum, i;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    for(i=1; i<=rank+1; i++)
        fact = fact * i;

    MPI_Reduce(&fact, &factsum, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
```

```

    if(rank==0)
        printf("Sum of all the factorial=%d",factsum);

    MPI_Finalize();
    exit(0);
}

```

Lab Exercises:

- 1) Write a MPI program using N processes to find $1! + 2! + \dots + N!$. Use scan. Also, handle different errors using error handling routines.
- 2) Write a MPI program to calculate π -value by integrating $f(x) = 4/(1+x^2)$. Area under the curve is divided into rectangles and the rectangles are distributed to the processors. Also handle different errors using error handling routines,
- 3) Write a MPI program to read a 3 X 3 matrix. Enter an element to be searched in the root process. Find the number of occurrences of this element in the matrix using three processes. Also, handle different errors using error handling routines.
- 4) Write a MPI program to read 4 X 4 matrix and display the following output using four processes.

I/p matrix: 1 2 3 4

O/p matrix: 1 2 3 4

1 2 3 1
1 1 1 1
2 1 2 1

2 4 6 5
3 5 7 6
5 6 9 7

Additional Exercises:

- 1) Write a MPI program to read a word of length N. Using N processes including the root get output word with the pattern as shown in example. Display the resultant output word in the root.
Example: Input : PCAP Output : PCCAAAPPPP
- 2) Write a MPI program to read matrix A of size 5 X 5. It produces a resultant matrix B of size 5 X 5. It sets all the principal diagonal elements of B matrix with 0. It replaces each row elements in the B matrix in the following manner: If the element is below the principal diagonal it replaces it with the maximum value of the column in the A matrix having the same row number of B. If the element is above the principal diagonal it replaces it with the minimum value of the column in the A matrix having the same row number of B. Produce the B Matrix using 5 processes. Use only Collective communication routines except broadcast routine.

Example:

A				
1	2	3	4	5
5	4	3	2	1
10	3	13	14	15
11	22	11	33	44
1	12	5	4	6

B				
0	1	1	1	1
22	0	2	2	2
13	13	0	3	3
33	33	33	0	2

OpenCL Introduction and programs on Vectors

Objectives:

In this lab, student will be able to

1. Understand how to write kernel code in OpenCL
2. Learn how to use different APIs in writing the host code in OpenCL
3. Study the execution environment necessary to execute OpenCL programs

I. OpenCL Introduction

- Open Computing Language is a heterogeneous programming framework that is managed by Khronos group.
- OpenCL is a framework for developing applications that execute across a range of device type made by different vendors.

OpenCL Specification: It is defined in four parts called Models.

Platform Model: Specifies that there is one processor coordinating the execution (the host) and one or more processors capable of executing OpenCL C code (the devices).

Execution Model: Defines how the **OpenCL environment** is configured on the host and how kernels are executed on the device. This includes

- setting up an **OpenCL context** on the host
- providing **mechanisms for host-device interaction**
- Defining a **concurrency model** used for execution on devices

Memory Model: Defines the abstract memory hierarchy that kernels use regardless of the actual underlying memory architecture.


Programming Model: Defines how the concurrency model is mapped to physical hardware

Platform and Devices

Defines the **roles of the host and devices** and provides an abstract hardware model for devices. Devices that a platform can target are limited to those with which a vendor knows how to interact. If company A's platform is chosen, it can't communicate with company B's GPU.

- Platform model defines a **device as an array of compute units**, with each compute unit functionally independent from the rest.
- These compute units are further divided into **processing elements**.

Host and Device Interaction



II. Different APIs used for writing OpenCL programs

Discover & initialize platforms

- CLGetPlatformIDs() – It is used to discover the set of available platforms for a given system.
- This function will be called twice by an application.
- The **first call** passes an unsigned int pointer as the num_platforms argument and NULL is passed as a platform's argument.
- The pointer is **populated with the available number of platforms**.
- Then the programmer can allocate space to hold the platform information.

```
cl_int clGetPlatformIDs ( cl_uint num_entries, cl_platform_id *platforms,  
                          cl_uint *num_platforms);
```

- For the second call, cl_platform_id pointer is passed to function with enough space allocated for num_entries platforms.

```
cl_int clGetPlatformIDs (cl_uint num_entries, cl_platform_id *platforms,  
                          cl_uint *num_platforms);
```

Discover & initialize Devices

cl_GetDeviceIDs() call works very similar to **cl_GetPlatformIDs()**. It takes the additional arguments of a **platform** and a **device type**.

Device type argument can be used to limit the devices to GPUs only (CL_DEVICE_TYPE_GPU), CPUs only (CL_DEVICE_TYPE_CPU), all devices (CL_DEVICE_TYPE_ALL)

```
cl_int clGetDeviceIDs ( cl_platform_id platform, cl_device_type device_type,  
                        cl_uint num_entries, cl_device_id *devices, cl_uint *num_devices);
```

```
cl_int clGetDeviceIDs (cl_platform_id platform, cl_device_type device_type, cl_uint  
                      num_entries, cl_device_id *devices, cl_uint *num_devices);
```

Execution Environment- Context

- Context is an abstract container that exists on the host.
- A context coordinates
 - the mechanisms for **host-device interactions**
 - **Manages the memory objects** that are available to the devices
 - **Keep track of the programs and kernels** that are created for each device

```
cl_context clCreateContext ( const cl_context_properties *properties,  
                           cl_uint num_devices, const cl_device_id *devices, void  
                           (CL_CALLBACK *pfn_notify)( const char *errinfo, const void  
                           *private_info, size_t cb, void *user_data), void *user_data,  
                           cl_int *errcode_ret);
```

Create command queue

- **Communication with a device** occurs by submitting commands to a command queue.
- Once the host decides which devices to work with and a context is created, **one command queue needs to be created per device**.
- Whenever the host needs an action to be performed by a device, it will submit commands to the proper command queue.

```
cl_command_queue clCreateCommandQueue( cl_context context,  
                                       cl_device_id device,  
                                       cl_command_queue_properties properties,  
                                       cl_int* errcode_ret);
```

- The **parameter** properties by default is **INORDER** ie command queue commands are pulled from the queue in the order they were received.

Memory objects

- In OpenCL the **data needs to be physically present on a device** before execution can begin. In order for data to be transferred to a device it must first be **encapsulated as a memory object**.

Two types of Memory objects available in OpenCL are:

- buffers
- images

Buffers

- Whenever a **memory object** is created, it is **valid** only **within a single context**.
- Creating a buffer requires supplying the size of the buffer and a context in which the buffer will be allocated.
- It is visible for **all devices associated with the context**.

Create device buffers

```
cl_mem clCreateBuffer ( cl_context context, cl_mem_flags flags, size_t size, void *host_ptr,
                      cl_int *errcode_ret);
```

flags can be:

CL_MEM_READ_ONLY or CL_MEM_WRITE_ONLY or CL_MEM_READ_WRITE

Write host data to device buffers

- Data contained in host memory is transferred to and from an OpenCL buffer using the commands **CLEnqueueWriteBuffer()** and **CLEnqueueReadBuffer()** respectively.

```
cl_int clEnqueueWriteBuffer ( cl_command_queue command_queue, cl_mem buffer,
                             cl_bool blocking_write, size_t offset, size_t cb, const void
                             *ptr, cl_uint num_events_in_wait_list, const cl_event
                             *event_wait_list, cl_event *event);
```

- **blocking_write** is set to **CL_TRUE** if the transfer into an OpenCL buffer should complete before function returns.
- **offset** The offset in bytes in the buffer object to write to.
- **cb** The size in bytes of data being written.
- **ptr** The pointer to buffer in host memory where data is to be written from

- If **event_wait_list** is NULL, then this particular command does not wait on any event to complete. If event_wait_list is NULL, num_events_in_wait_list must be 0.
- If event_wait_list is not NULL, the list of events pointed to by event_wait_list must be valid and num_events_in_wait_list must be greater than 0.
- **event** Returns an event object that **identifies this particular write command** and can be used to query or queue a wait for this particular command to complete.

Creating a OpenCL program Object

- OpenCL C (kernel) code is called a program
- OpenCL programs are compiled at runtime through a series of API calls

The process of creating a kernel is as follows:

- The OpenCL C source code (kernel) is stored in a character stream/string.
- The source code is turned into a **program object** cl_program by calling **clCreateProgramWithSource()**
- The program object is then compiled for one of OpenCL devices with **clBuildProgram()**

Kernels & OpenCL execution Model

- **Kernels** are the part of OpenCL program that actually executed on a device. Return type must be void

```
__kernel void vecadd(__global int *C, __global int* A, __global int *B)
{
    int tid = get_global_id(0); // OpenCL intrinsic function
    C[tid] = A[tid] + B[tid];
}
```

- The unit of concurrent execution in OpenCL is a **work item**.
- Single iteration of a loop is mapped to a work-item.
- We tell OpenCL runtime to generate as many work-items as elements in the input and output arrays and allow the runtime to map those work-items to the underlying hardware.
- **get_global_id()** will give the **position of the current work item**.

Transfer kernel code to a character array

```
char *source_str;
size_t source_size;
fp = fopen("Hello.cl", "r");
if (!fp) {
    fprintf(stderr, "Failed to load kernel.\n");
}
```



```

    getchar();
    exit(1);
}
source_str = (char*)malloc(MAX_SOURCE_SIZE);
source_size = fread( source_str, 1, MAX_SOURCE_SIZE, fp);
fclose( fp );

```

Create and compile the program

```

cl_program    clCreateProgramWithSource(  cl_context    context,  cl_uint    count,
const char**  strings, const size_t *lengths, cl_int *errcode_ret);

```

```

cl_int clBuildProgram( cl_program program, cl_uint num_devices, cl_device_id *
                      device_list, const char *options, void (ptr_notify *) (cl_program,
                      void *user_data), void *user_data);

```

If device_list is **NULL** value, the program executable is built for all devices associated with program for which a source or binary has been loaded.

If device_list is a **non-NULL** value, the program executable is built for devices specified in this list for which a source or binary has been loaded.

Create the Kernel

- The final stage to obtain a CL kernel object that can be used to execute kernels on a device is to extract the kernel from the **program object** (cl_program).
- The name of the kernel is passed to **clCreateKernel** function along with the program object.
- The kernel object will be returned if the program object was valid and particular kernel is formed.

```

cl_kernel clCreateKernel( cl_program program, const char* kernel_name ,
                        cl_int * errcode_ret);

```

- kernel_name - A function name in the program declared with the __kernel qualifier

Set the kernel arguments

```

cl_int clSetKernelArg( cl_kernel kernel, cl_uint arg_index , size_t
                      arg_size, const void * arg_value);

```

- Arg_index-Arguments to the kernel are referred by indices that go from 0 for the leftmost argument to n - 1, where n is the total number of arguments declared by a kernel.
- We must specify each kernel argument individually using the function **clSetKernelArg()**.
- When the kernel is executed this information is used to transfer arguments to the device.

OpenCL execution Model

- When a kernel is executed, the programmer specifies the number of work items that should be created as an N dimensional range, ie NDRange. NDRange is 1 or 2 or 3 dimensional index space of work items that will often map to the dimension of either the input or output data.

Set global & local work size

- Scalability comes from dividing the work items of an NDRange into smaller equally sized workgroups using the same N dimensions.
- For example Work Dimension is 1 and if we set
size_t global_work_size = 1024;

size_t local_work_size=128;

Since there are 1024 total work-items and 128 work-items / work-group, a simple division of $1024 / 128 = 8$ **work-groups** are created.

- The global_work_size is the **total number of work-items (WI)**
- The local_work_size is the **number of work-items per work-group (WI/WG)**
- The number of work-groups is **the global_work_size / local_work_size or WG**

Enqueue the kernel for execution:

```
cl_int clEnqueueNDRangeKernel( cl_command_queue command_queue, cl_kernel kernel,
                               cl_uint work_dim, const size_t *global_work_offset, const
                               size_t *global_work_size, const size_t *local_work_size,
                               cl_uint num_events_in_wait_list, const cl_event
                               *event_wait_list, cl_event *event);
```

4 fields are related to work-item creation:

- **work_dim:** specifies the number of dimensions in which work-item will be created
- **global_work_size:** specifies the number of work items in each dimension of the ND range
- **local_work_size:** specifies the number of work items in each dimension of the workgroup
- **global_work_offset:** used to provide global IDs to the work items that do not start from zero

Read the output buffer back to the host

```
cl_int clEnqueueReadBuffer ( cl_command_queue command_queue, cl_mem buffer,
                             cl_bool blocking_write, size_t offset, size_t cb, const
                             void *ptr, cl_uint num_events_in_wait_list, const
                             cl_event *event_wait_list, cl_event *event);
```

Barrier Operations for a command queue

- Two types of barrier operations for a command queue:

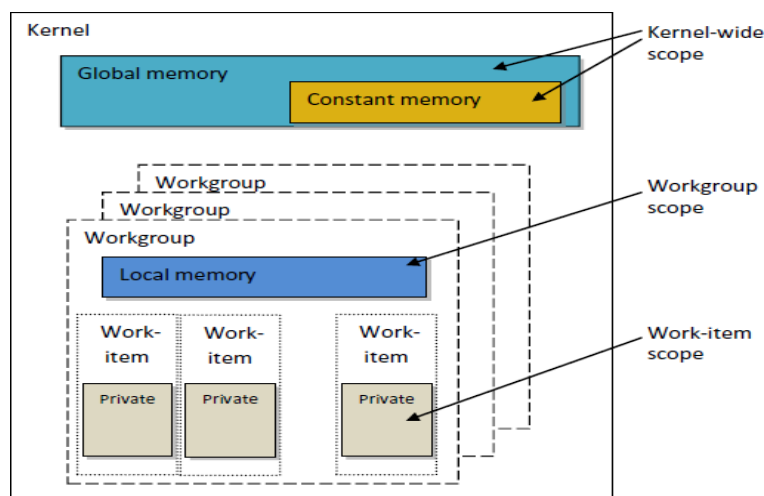
```
cl_int clFinish (cl_command_queue cmdQueue);
```

This function blocks until all of the commands in a command queue have completely executed.

```
cl_int clFlush (cl_command_queue cmdQueue);
```

This function blocks until all of the commands in a command queue have been removed from the command queue.

The abstract memory model



- Global Memory(_global)** – Is **visible to all compute units on the device**, similar to main memory on a CPU based system.
- Whenever data is transferred from host to device and transferred back from device to host, it must reside in global memory.
- Constant Memory(_constant)** – It is modeled as a part of global memory, so memory objects that are transferred to global memory can be specified as constant variables whose values never change fall into this category.
- Local Memory (_local)** – It is modeled as being **shared by a workgroup**.

- **Private Memory** – It is **unique to an individual work item**. Local variables and non_pointer kernel arguments are private by default.

Release OpenCL resources:

```
// Free OpenCL resources
clReleaseKernel(kernel);
clReleaseProgram(program);
clReleaseCommandQueue(cmdQueue);
clReleaseMemObject(bufferA);
clReleaseMemObject(bufferB);
clReleaseMemObject(bufferC);
clReleaseContext(context);
// Free host resources
free(A); //arrays
free(B);
free(C);
free(platforms);
free(devices);
```

Solved Exercise:

Write an OpenCL program to read two arrays A and B of same size N. Find the sum of corresponding array elements. Store the result in array C.

```
// A COMPLETE PROGRAM FOR VECTOR-VECTOR ADDITION

#include <stdio.h>
#include <CL/cl.h>
#include <stdlib.h>

//Max source size of the kernel string
#define MAX_SOURCE_SIZE (0x100000)

int main(void)
{
    // Create the two input vectors
    int i;
    int LIST_SIZE;
    printf("Enter how many elements:");
    scanf("%d",&LIST_SIZE);
```

```

int *A = (int*) malloc (sizeof (int) * LIST_SIZE);

//Initialize the input vectors
for(i = 0; i < LIST_SIZE; i++)
{
    A[i] = i; //if LIST_SIZE is very large
}

int *B = (int*)malloc(sizeof(int)*LIST_SIZE);
//Initialize the input vectors
for(i = 0; i < LIST_SIZE; i++)
{
    B[i] = i+10;
}

// Load the kernel source code into the array source_str
FILE *fp;
char *source_str;
size_t source_size;

fp = fopen("vectorCLKernel.cl", "r");

if (!fp)
{
    fprintf(stderr, "Failed to load kernel.\n");
    getchar();
    exit(1);
}
source_str = (char*)malloc(MAX_SOURCE_SIZE);
source_size = fread( source_str, 1, MAX_SOURCE_SIZE, fp);

fclose( fp );

// Get platform and device information
cl_platform_id platform_id = NULL;
cl_device_id device_id = NULL;
cl_uint ret_num_devices;
cl_uint ret_num_platforms;

cl_int ret = clGetPlatformIDs(1, &platform_id, &ret_num_platforms);
ret = clGetDeviceIDs( platform_id, CL_DEVICE_TYPE_GPU, 1,&device_id,
&ret_num_devices);

```

```

// Create an OpenCL context
cl_context context = clCreateContext( NULL, 1, &device_id, NULL, NULL, &ret);

// Create a command queue
cl_command_queue command_queue = clCreateCommandQueue(context, device_id,
NULL, &ret);

// Create memory buffers on the device for each vector A, B and C
cl_mem a_mem_obj = clCreateBuffer(context, CL_MEM_READ_ONLY, LIST_SIZE *
sizeof(int), NULL, &ret);
cl_mem b_mem_obj = clCreateBuffer(context, CL_MEM_READ_ONLY, LIST_SIZE *
sizeof(int), NULL, &ret);

cl_mem c_mem_obj = clCreateBuffer(context, CL_MEM_WRITE_ONLY, LIST_SIZE *
sizeof(int), NULL, &ret);

// Copy the lists A and B to their respective memory buffers
ret = clEnqueueWriteBuffer(command_queue, a_mem_obj, CL_TRUE, 0, LIST_SIZE *
sizeof(int), A, 0, NULL, NULL);

ret = clEnqueueWriteBuffer(command_queue, b_mem_obj, CL_TRUE, 0, LIST_SIZE *
sizeof(int), B, 0, NULL, NULL);

// Create a program from the kernel source
cl_program program = clCreateProgramWithSource(context, 1, (const
char**)&source_str, (const size_t *)&source_size, &ret);

// Build the program
ret = clBuildProgram(program, 1, &device_id, NULL, NULL, NULL);

// Create the OpenCL kernel object
cl_kernel kernel = clCreateKernel(program, "vector_add", &ret);

// Set the arguments of the kernel
ret = clSetKernelArg(kernel, 0, sizeof(cl_mem), (void *)&a_mem_obj);
ret = clSetKernelArg(kernel, 1, sizeof(cl_mem), (void *)&b_mem_obj);
ret = clSetKernelArg(kernel, 2, sizeof(cl_mem), (void *)&c_mem_obj);

// Execute the OpenCL kernel on the array
size_t global_item_size = LIST_SIZE;
size_t local_item_size = 1;

//Execute the kernel on the device
cl_event event;
ret = clEnqueueNDRangeKernel(command_queue, kernel, 1, NULL, &global_item_size,
&local_item_size, 0, NULL, NULL);

```

```

ret = clFinish(command_queue);

    // Read the memory buffer C on the device to the local variable C
    int *C = (int*)malloc(sizeof(int)*LIST_SIZE);
    ret = clEnqueueReadBuffer(command_queue, c_mem_obj, CL_TRUE, 0,LIST_SIZE *
    sizeof(int), C, 0, NULL, NULL);

    // Display the result to the screen
    for(i = 0; i < LIST_SIZE; i++)
        printf("%d + %d = %d\n", A[i], B[i], C[i]);

    // Clean up
    ret = clFlush(command_queue);
    ret = clReleaseKernel(kernel);
    ret = clReleaseProgram(program);
    ret = clReleaseMemObject(a_mem_obj);
    ret =
    clReleaseMemObject(b_mem_obj); ret =
    clReleaseMemObject(c_mem_obj);
    ret =
    clReleaseCommandQueue(command_queue); ret =
    clReleaseContext(context);

    free(A);
    free(B);
    free(C);
    getchar();
    return 0;
}

// vectorCLKernel.cl
__kernel void vector_add(__global int *A, __global int *B,__global int *C)
{
    // Get the index of the current work item
    int i = get_global_id(0);
    // Do the operation
    C[i] = A[i] + B[i];
}

```

Steps to execute an OpenCL program is provided in the form of video which is made available in individual systems. However, the basic installation steps are given as follows.

// To install OpenCL in Ubuntu OS (Intel GPU)

\$sudo apt-get update

\$sudo apt-get install ocl-icd-opencl-dev

// To edit a file, create a file in any text editor such as vim or gedit with an extension .cl

//To Compile a OpenCL program

\$ gcc filename.cl -o filename.out -lOpenCL

//To Run a OpenCL program

\$./filename.out

Lab Exercises:

- 1) Write an OpenCL program which takes N integers as input. It converts these integers into their corresponding octal values and stores the result in another array in parallel.
- 2) Write an OpenCL program which takes N binary numbers as input and stores the one's complement of each element in another array in parallel.
- 3) Write an OpenCL program which takes N (Even No.) integers into an 1D array and modifies the same array by swapping adjacent elements in parallel.
- 4) Write an OpenCL program which takes N binary numbers as input. It converts these values into their corresponding decimal values(base 10) and stores the result in another array in parallel.

Additional Exercises:

- 1) Write an OpenCL program which reads an array A with N number of integers and checks each element is prime or not. If prime, store the number as it is else store the square of the number in another array B.
- 2) A source Router is connected to M other routers. Each of the M router has N devices. Each device has unique IP address(version 4) with subnet mask 255.255.255.0 Assume the source router has a forwarding table with M rows and N columns containing last octet of the corresponding IP address. Assume a new packet has arrived at the source router, determine using a OpenCL parallel programming to which router, the source router should forward the packet so that it reaches the correct destination? Write the kernel, kernel invocation and display of result section of the program.

OpenCL programs on Strings, Sorting **and to Check the execution time in OpenCL**

Objectives:

In this lab, student will be able to

1. Understand how to write kernel code in OpenCL to perform operations on strings
2. Learn how to use different APIs in writing the host code in OpenCL in order to deal with strings
3. Learn different APIs used in OpenCL to check the execution time taken by kernel

APIs used to find the time taken by kernel

- Execution time taken by the kernel can be extracted by using the function **clGetEventProfilingInfo()** after the execution of the kernel by the function **clEnqueueNDRangeKernel()**.
- **clFinish()** function must be used after **clEnqueueNDRangeKernel()** so that finish time is calculated after the execution of kernel.

```
cl_int clGetEventProfilingInfo ( cl_event event, cl_profiling_info param_name,  
                                size_t param_value_size, void *param_value, size_t  
                                *param_value_size_ret);
```

- **cl_profiling_info** can have following types:

CL_PROFILING_COMMAND_START: (cl_ulong): A 64-bit value that describes the current device time counter in nanoseconds when the command identified by event starts execution on the device.

CL_PROFILING_COMMAND_END: (cl_ulong): A 64-bit value that describes the current device time counter in nanoseconds when the command identified by event has finished execution on the device.

- Profiling of OpenCL commands can be enabled by using a command-queue created with **CL_QUEUE_PROFILING_ENABLE** flag set in properties argument to **clCreateCommandQueue()**.

Solved Exercise:

Write an OpenCL program which reads a string as input and it toggles every character of this string in parallel. Store the resultant string in another character array. Also find the execution time taken by kernel.

```
#include <stdio.h>
#include<time.h>
#include <CL/cl.h>
#include<string.h>
#include<stdlib.h>
#include<conio.h>
#define MAX_SOURCE_SIZE (0x100000)

int main(void)
{
    time_t start, end;
    start = clock();

    char tempstr[20486];

    //Initialize the input string
    int i;

    for(i=0; i<20485; i++)
        tempstr[i]='A';
    tempstr[20485]='\0';

    int len= strlen(tempstr);
    len++;

    char *str = (char*)malloc(sizeof(char)*len);

    strcpy(str,tempstr);

    // Load the kernel source code into the array source_str
    FILE *fp;
    char *source_str;
    size_t source_size;

    fp = fopen("strtoggle.cl", "r");
    if (!fp) {
        fprintf(stderr, "Failed to load kernel.\n");
        getchar();
        exit(1);
    }
}
```

```

source_str = (char*)malloc(MAX_SOURCE_SIZE);
source_size = fread( source_str, 1, MAX_SOURCE_SIZE, fp);
fclose( fp );

cl_platform_id platform_id = NULL;
cl_device_id device_id = NULL;
cl_uint ret_num_devices;
cl_uint ret_num_platforms;

cl_int ret = clGetPlatformIDs(1, &platform_id, &ret_num_platforms);

ret = clGetDeviceIDs( platform_id, CL_DEVICE_TYPE_GPU, 1,&device_id,
&ret_num_devices);

// Create an OpenCL context
cl_context context = clCreateContext( NULL, 1, &device_id, NULL, NULL, &ret);

// Create a command queue
cl_command_queue command_queue = clCreateCommandQueue(context,
device_id, CL_QUEUE_PROFILING_ENABLE, &ret);

// Create memory buffers on the device for input and output string
cl_mem s_mem_obj = clCreateBuffer(context, CL_MEM_READ_ONLY, len*
sizeof(char), NULL, &ret);

cl_mem t_mem_obj = clCreateBuffer(context, CL_MEM_WRITE_ONLY, len*
sizeof(char), NULL, &ret);

// Copy the input string into respective memory buffer
ret = clEnqueueWriteBuffer(command_queue, s_mem_obj, CL_TRUE, 0, len *
sizeof(char), str, 0, NULL, NULL);

// Create a program from the kernel source
cl_program program = clCreateProgramWithSource(context, 1,(const char
**)&source_str, (const size_t *)&source_size, &ret);

// Build the program
ret = clBuildProgram(program, 1, &device_id, NULL, NULL, NULL);

// Create the OpenCL kernel
cl_kernel kernel = clCreateKernel(program, "str_chgcase", &ret);

```

```

// Set the arguments of the kernel
ret = clSetKernelArg(kernel, 0, sizeof(cl_mem), (void *)&s_mem_obj);
ret = clSetKernelArg(kernel, 1, sizeof(cl_mem), (void *)&t_mem_obj);

// Set the global work size as string length
size_t global_item_size = len; // Process the entire lists
size_t local_item_size = 1;

//Execute the OpenCL kernel for entire string in parallel
cl_event event;
ret = clEnqueueNDRangeKernel(command_queue, kernel, 1, NULL,
&global_item_size, &local_item_size, 0, NULL, &event);
time_t stime=clock();

//kernel execution must be finished before calculating time
ret = clFinish(command_queue);

cl_ulong time_start, time_end;
double total_time;

//Find the kernel execution start time
clGetEventProfilingInfo(event, CL_PROFILING_COMMAND_START,
sizeof(time_start), &time_start, NULL);

//Find the kernel execution end time
clGetEventProfilingInfo(event, CL_PROFILING_COMMAND_END,
sizeof(time_end), &time_end, NULL);
total_time = double (time_end - time_start);

// Read the result in memory buffer on the device to the local variable strres
char *strres = (char*)malloc(sizeof(char)*len);

ret = clEnqueueReadBuffer(command_queue, t_mem_obj, CL_TRUE, 0, len*
sizeof(char), strres, 0, NULL, NULL);
printf("\nDone");
strres[len-1]='\0';
printf("\nResultant toggled string :%s",strres);
getchar();

ret = clReleaseKernel(kernel);
ret = clReleaseProgram(program);
ret = clReleaseMemObject(s_mem_obj);
ret = clReleaseMemObject(t_mem_obj);
ret = clReleaseCommandQueue(command_queue);
ret = clReleaseContext(context);

```

```

        end = clock();

        printf("\n\n Time taken to execute the KERNEL in milliseconds = %0.3f
msec\n\n", total_time/1000000);
        printf( "\nTime taken to execute the whole program in seconds: %0.3f
Seconds\n", (end-start)/(double)CLOCKS_PER_SEC );

free(str);
free(strres);
getch();
return 0;
}

// strtoggle.cl
__kernel void str_chgcase (__global char *A, __global char *B)
{

    int i = get_global_id(0);
    if(A[i] >='A' && A[i] <='Z')
        B[i] = A[i] + 32;
    else
        B[i] = A[i] - 32;

}

```

Lab Exercises:

- 1) Write an OpenCL program that takes a string S as input and one integer value N. The program produces the output string N times as shown below in parallel:
I/p: S = Hello N = 3
O/p String: HelloHelloHello
Note: Each work item copies entire S
- 2) Write an OpenCL program to read an array of N integer values. Sort each element of this array in parallel using parallel selection sort and store the result in another array.
- 3) Write an OpenCL program to read an integer array of size N. Sort this array using odd-even transposition sorting. Use 2 kernels.
- 4) Write an OpenCL program which reads a string consisting of N words and reverse each word of it in parallel. Use N number of processes.

Additional Exercises:

- 1) Write an OpenCL program that takes a string S as input and one integer value N. Produces string N times as follows in parallel:

I/p: S = Hello N = 3

O/p String: HelloHelloHello

Note : Each work item copies same character from the Input N times to the required position)

- 2) Write an OpenCL program which reads a string S and produces output string T as follows:

I/p: S: Hai O/p : T: Haaiii

Note: Each work item stores a character from input string S required number of times in T.

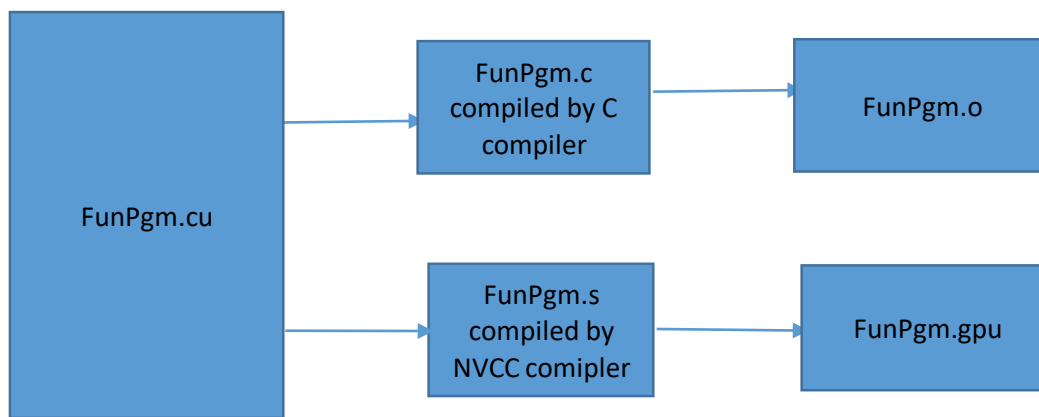
Programs on Arrays in CUDA

Objectives:

In this lab, student will be able to

1. Know the basics of Computing Unified Device Architecture (CUDA).
2. Learn program structure of CUDA.
3. Learn about CUDA 1D blocks and threads
4. Write simple programs on one dimensional arrays
5. Learn mathematical functions in CUDA

About CUDA: CUDA is a platform for performing massively parallel computations on graphics accelerators. CUDA was developed by NVIDIA. It was first available with their G8X line of graphics cards. CUDA presents a unique opportunity to develop widely-deployed parallel applications. The CUDA programs are compiled as follows.



FunPgm.cu is compiled by both C compiler and Nvidia CUDA C compiler (NVCC compiler). If you have both main.c and Funpgm.cu then you can call cuda API's in main.c but keep in mind that you cannot call kernel from main.c. To call the kernel file extension must be .cu.

As in OpenCL CPU is the host and its memory the host memory and GPU is the device and its memory device memory. Serial code will be run on host and parallel code will be run on device.

1. Copy data from host memory to device memory.
2. Load device program and execute, caching data on chip for performance.
3. Copy result from device memory to host memory.

CUDA threads, blocks and grid

Thread – Distributed by the CUDA runtime. A single path of execution there can be multiple threads in a program.


(identified by threadIdx)

CUDA Thread




Block – A user defined group of 1 to 512 threads.

(identified by blockIdx)



CUDA Block

Grid – A group of one or more blocks. A grid is created for each CUDA kernel function



CUDA
GRID

Some of the calculations for indexing the thread is given below.

1D grid of 1D blocks

__device__

```
int getGlobalIdx_1D_1D(){  
    return blockIdx.x * blockDim.x + threadIdx.x;  
}
```

1D grid of 2D blocks

__device__

```
int getGlobalIdx_1D_2D(){  
    return blockIdx.x * blockDim.x * blockDim.y  
        + threadIdx.y * blockDim.x + threadIdx.x;  
}
```

1D grid of 3D blocks

__device__

```
int getGlobalIdx_1D_3D(){  
    return blockIdx.x * blockDim.x * blockDim.y * blockDim.z  
        + threadIdx.z * blockDim.y * blockDim.x  
        + threadIdx.y * blockDim.x + threadIdx.x;  
}
```

2D grid of 1D blocks

```
__device__ int getGlobalIdx_2D_1D(){  
    int blockId = blockIdx.y * gridDim.x + blockIdx.x;  
    int threadId = blockId * blockDim.x + threadIdx.x;
```

```
return threadIdx;
```

```
}
```

2D grid of 2D blocks

```
__device__
```

```
int getGlobalIdx_2D_2D(){
```

```
int blockId = blockIdx.x + blockIdx.y * gridDim.x;
```

```
int threadId = blockId * (blockDim.x * blockDim.y)
```

```
+ (threadIdx.y * blockDim.x) + threadIdx.x;
```

```
return threadId;
```

```
}
```

2D grid of 3D blocks

```
__device__
```

```
int getGlobalIdx_2D_3D(){
```

```
int blockId = blockIdx.x + blockIdx.y * gridDim.x;
```

```
int threadId = blockId * (blockDim.x * blockDim.y * blockDim.z)
```

```
+ (threadIdx.z * (blockDim.x * blockDim.y))
```

```
+ (threadIdx.y * blockDim.x) + threadIdx.x;
```

```
return threadId;
```

```
}
```

3D grid of 1D blocks

```
__device__
```

```
int getGlobalIdx_3D_1D(){
```

```
int blockId = blockIdx.x + blockIdx.y * gridDim.x
```

```
+ gridDim.x * gridDim.y * blockIdx.z;
```

```
int threadId = blockId * blockDim.x + threadIdx.x;
```

```

return threadIdx;

}

3D grid of 2D blocks

__device__

int getGlobalIdx_3D_2D(){

int blockIdx = blockIdx.x + blockIdx.y * gridDim.x

+ gridDim.x * gridDim.y * blockIdx.z;

int threadIdx = blockIdx * (blockDim.x * blockDim.y)

+ (threadIdx.y * blockDim.x) + threadIdx.x;

return threadIdx;

}

3D grid of 3D blocks

__device__

int getGlobalIdx_3D_3D(){

int blockIdx = blockIdx.x + blockIdx.y * gridDim.x

+ gridDim.x * gridDim.y * blockIdx.z;

int threadIdx = blockIdx * (blockDim.x * blockDim.y * blockDim.z)

+ (threadIdx.z * (blockDim.x * blockDim.y))

+ (threadIdx.y * blockDim.x) + threadIdx.x;

return threadIdx;

}

```

Solved Exercise: Program to add two numbers.

```
#include "cuda_runtime.h"
#include "device_launch_parameters.h"

__global__ void add(int *a, int *b, int *c)
{
    *c = *a + *b;
}

int main(void) {
    int a, b, c;           // host copies of variables a, b & c
    int *d_a, *d_b, *d_c; // device copies of variables a, b & c
    int size = sizeof(int);

    // Allocate space for device copies of a, b, c
    cudaMalloc((void **)&d_a, size);
    cudaMalloc((void **)&d_b, size);
    cudaMalloc((void **)&d_c, size);
    // Setup input values
    a = 3;
    b = 5;
    // Copy inputs to device
    cudaMemcpy(d_a, &a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, &b, size, cudaMemcpyHostToDevice);
    // Launch add() kernel on GPU
    add<<<1,1>>>>(d_a, d_b, d_c);
    // Copy result back to host
    cudaMemcpy(&c, d_c, size, cudaMemcpyDeviceToHost);
    printf("Result : %d",c);
    // Cleanup
    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);
    return 0;
}
```

Explanation:

add is the function which runs on device.

```
cudaMalloc ((void **)&d_a, size);
```

cudaMalloc will allocate memory of size bytes given as second argument to variable passed as first argument.

```
cudaMemcpy (Destination,Source,Size,Direction);  
  
cudaMemcpy (d_a, &a, size, cudaMemcpyHostToDevice);  
  
cudaMemcpy (&c, d_c, size, cudaMemcpyDeviceToHost);
```

cudaMemcpy copies the variables from host to device or device to host based on the direction which is either cudaMemcpyHostToDevice or cudaMemcpyDeviceToHost. Value of size bytes long is copied from source to destination.

cudaFree frees the memory allocated by cudaMalloc.

The add function is called like this add<<<1,1>>>(d_a,d_b,d_c). The add is followed by three angular brackets then the number of blocks, threads per block then corresponding closing angular brackets then how many arguments the function add takes is enclosed within parenthesis. If you want to add N elements you can achieve it in two ways either having N blocks or having N threads.

That is pass an array with following function calls

add<<< N,1>>> (d_a,d_b,d_c) or add<<< 1, N>>> (d_a,d_b,d_c)

Few Mathematical functions in CUDA:

Major Single-Precision floating point functions: Single precision functions work on float value(32 bit). A float value is stored in IEEE 754 format.

Function	Description
sqrtf(x)	Square root function
expf(x)	Exponentiation function. Base = e
exp2f(x)	Exponentiation function. Base = 2
exp10f(x)	Exponentiation function. Base = 10
logf(x)	Logarithmic function. Base=e
log2f(x)	Logarithmic function. Base=2
log10f(x)	Logarithmic function. Base=10
sinf(x)	sine function

cosf(x)	cos function
tanf(x)	tan function
powf(x,y)	power function
truncf(x)	truncation function
roundf(x)	round function
ceilf(x)	ceil function
floorf(x)	floor function

Major Double-Precision floating point functions: Double precision functions work on double value(64 bit). A double value is stored in IEEE 754 format.

Function	Description
sqrt(x)	Square root function
exp(x)	Exponentiation function. Base = e
exp2(x)	Exponentiation function. Base = 2
exp10(x)	Exponentiation function. Base = 10
log(x)	Logarithmic function. Base=e
log2(x)	Logarithmic function. Base=2
log10(x)	Logarithmic function. Base=10
sin(x)	sine function
cos(x)	cos function
tan(x)	tan function
pow(x,y)	power function
trunc(x)	truncation function
round(x)	round function

ceil(x)	ceil function
floor(x)	floor function

Steps to execute a CUDA program is provided in the form of video which is made available in individual systems.

Lab Exercises:

1. Write a program in CUDA to add two vectors of length N using
a) block size as N b) N threads
2. Implement a CUDA program to add two vectors of length N by keeping the number of threads per block as 256 (constant) and vary the number of blocks to handle N elements.
3. Write a program in CUDA which performs convolution operation on one dimensional input array N of size *width* using a mask array M of size *mask_width* to produce the resultant one dimensional array P of size *width*.
4. Write a program in CUDA to process a 1D array containing angles in radians to generate sine of the angles in the output array. Use appropriate function.

Additional Exercises:

1. Write a program in CUDA to perform linear algebra function of the form $y = \alpha x + y$, where x and y are vectors and α is a scalar value.
2. Write a program in CUDA to sort every row of a matrix using selection sort.
3. Write a program in CUDA to perform odd even transposition sort in parallel.

Programs on Strings in CUDA

Objectives:

In this lab, student will be able to

1. Write simple programs on Strings
2. Learn to compute time of kernel execution
3. Learn about atomic functions
4. Learn to handle errors in the kernel

Arithmetic functions:

In a multithreaded scenario, the issue of data inconsistency will arise, if multiple threads modify a single shared memory variable. To overcome this atomic functions need to be used. List of atomic functions, their syntax and explanation is provided below.

atomicAdd():

```
int atomicAdd (int* address, int val);  
  
unsigned int atomicAdd(unsigned int* address, unsigned int val);  
  
float atomicAdd(float* address, float val);  
  
double atomicAdd(double* address, double val);
```

Reads the 16-bit, 32-bit or 64-bit word old located at the address address in global or shared memory, computes (old + val), and stores the result back to memory at the same address. These three operations are performed in one atomic transaction. The function returns old.

atomicSub():

```
int atomicSub(int* address, int val);  
  
unsigned int atomicSub(unsigned int* address, unsigned int val);
```

Reads the 32-bit word old located at the address address in global or shared memory, computes (old - val), and stores the result back to memory at the same address. These three operations are performed in one atomic transaction. The function returns old.

atomicExch():

```
int atomicExch(int* address, int val);  
  
unsigned int atomicExch(unsigned int* address, unsigned int val);  
  
float atomicExch(float* address, float val);
```

Reads the 32-bit word old located at the address address in global or shared memory and stores val back to memory at the same address. These two operations are performed in one atomic transaction. The function returns old.

atomicMin():

```
int atomicMin(int* address, int val);  
  
unsigned int atomicMin(unsigned int* address, unsigned int val);
```

Reads the 32-bit word old located at the address address in global or shared memory, computes the minimum of old and val, and stores the result back to memory at the same address. These three operations are performed in one atomic transaction. The function returns old.

atomicMax():

```
int atomicMax(int* address, int val);  
  
unsigned int atomicMax(unsigned int* address, unsigned int val);
```

Reads the 32-bit word old located at the address address in global or shared memory, computes the maximum of old and val, and stores the result back to memory at the same address. These three operations are performed in one atomic transaction. The function returns old.

atomicInc():

```
unsigned int atomicInc(unsigned int* address, unsigned int val);
```

Reads the 32-bit word old located at the address address in global or shared memory, computes $((old \geq val) ? 0 : (old+1))$, and stores the result back to memory at the same address. These three operations are performed in one atomic transaction. The function returns old.

atomicDec():

```
unsigned int atomicDec(unsigned int* address, unsigned int val);
```

Reads the 32-bit word old located at the address address in global or shared memory, computes $((old == 0) \vee (old > val)) ? val : (old - 1)$, and stores the result back to memory at the same address. These three operations are performed in one atomic transaction. The function returns old.

Solved Example:

A CUDA program which takes a string as input and determines the number of occurrences of a character 'a' in the string. This program uses atomicAdd() function.

```
#include "cuda_runtime.h"
#include "device_launch_parameters.h"

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <conio.h>
#define N 1024

__global__ void CUDACount(char* A, unsigned int *d_count){
    int i = threadIdx.x;
    if(A[i]=='a')
        atomicAdd(d_count,1);
}

int main() {
    char A[N];
    char *d_A;
    unsigned int *count=0,*d_count,*result;
    printf("Enter a string");
    gets(A);
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);
    cudaEventRecord(start, 0);
    cudaMalloc((void**)&d_A, strlen(A)*sizeof(char));
    cudaMalloc((void **)&d_count,sizeof(unsigned int));
    cudaMemcpy(d_A, A, strlen(A)*sizeof(char), cudaMemcpyHostToDevice);
    cudaMemcpy(d_count,count,sizeof(unsigned int),cudaMemcpyHostToDevice);
```

```

cudaError_t error = cudaGetLastError();
if (error != cudaSuccess) {
        printf("CUDA Error1: %s\n", cudaGetErrorString(error));
    }
    CUDACount<<<1,strlen(A)>>>(d_A,d_count);
error = cudaGetLastError();
if (error != cudaSuccess) {
        printf("CUDA Error2: %s\n", cudaGetErrorString(error));
    }

cudaEventRecord(stop, 0);
cudaEventSynchronize(stop);
float elapsedTime;
cudaEventElapsedTime(&elapsedTime, start, stop);
    cudaMemcpy(result, d_count, sizeof(unsigned int), cudaMemcpyDeviceToHost);
    printf("Total occurrences of a=%u",result);
printf("Time Taken=%f",elapsedTime);
    cudaFree(d_A);
    cudaFree(d_count);
    printf("\n");
    getch();
    return 0;
}

```

Explanation:

The kernel uses atomicAdd function with 1 as the value each time a character 'a' occurs in the string. The instructions given in bold are present to find the time. As in OpenCL you need to declare an event, register the event and record the time before kernel execution and after kernel execution. You have to synchronize the event so that main thread can capture the time of execution of kernel. After that **cudaEventElapsedTime(&elapsedTime, start, stop)** will give the difference between the recorded stop and start time and store the value in the variable elapsedTime which is of type float. A negative time value means there is something wrong in the CUDA code. To find it out you use the code given in bold and italics. It will display the error message present in CUDA code. Call it once before calling the kernel and once after calling kernel. If first call throws an error message then error is present in CUDA API which precedes the kernel. If second call throws the error message then error is present in the kernel code. .

Lab Exercises:

1. Write a program in CUDA to count the number of times a given word is repeated in a sentence. (Use Atomic function)

2. Write a CUDA program that reads a string *S* and produces the string *RS* as follows:

Input string *S*: PCAP Output string *RS*: PCAPPCAPCP

Note: Each work item copies required number of characters from *S* in *RS*.

Additional Exercises:

1) Write a CUDA program which reads a string consisting of N words and reverse each word of it in parallel.

2) Write a CUDA program that takes a string *Sin* as input and one integer value N and produces an output string , *Sout*, in parallel by concatenating input string *Sin*, N times as shown below.

Input: *Sin* = "Hello" N = 3

Output: *Sout* = "HelloHelloHello"

Note: Every thread copies the same character from the Input string *S*, N times to the required position.

3) Write a CUDA program which reads a string *Sin* and produces an output string *T* as shown below.

Input: *Sin*: "Hai"

Output: *T*: "Haaiii "

Note: Every thread stores a character from input string *Sin*, required number of times into output string *T*.

Programs on Matrix using CUDA

Objectives:

In this lab, student will be able to

1. Understand how to write kernel code in CUDA to perform operations on matrix
2. Learn about CUDA 2D blocks/threads
3. Write simple programs on two dimensional arrays

Solved Exercise:

Write a program in CUDA to find transpose of a matrix in parallel.

```
#include "cuda_runtime.h"
#include "device_launch_parameters.h"
#include <stdio.h>
#include <stdlib.h>

__global__ void transpose(int *a, int *t)
{
    int n=threadIdx.x, m=blockIdx.x, size=blockDim.x, size1=gridDim.x;
    t[n*size1+m]=a[m*size+n];
}

int main(void)
{
    int *a,*t, m,n,i,j;
    int *d_a,*d_t;
    printf("Enter the value of m: ");scanf("%d",&m);
    printf("Enter the value of n: ");scanf("%d",&n);
    int size=sizeof(int)*m*n;
    a=(int*)malloc(m*n*sizeof(int));
    c=(int*)malloc(m*n*sizeof(int));
    printf("Enter input matrix:\n");
    for(i=0;i<m*n;i++)
        scanf("%d",&a[i]);

    cudaMalloc((void**)&d_a,size);
    cudaMalloc((void**)&d_t,size);

    cudaMemcpy(d_a,a,size,cudaMemcpyHostToDevice);
```

```

transpose<<<m,n>>>(d_a,d_t);
cudaMemcpy(t,d_t,size,cudaMemcpyDeviceToHost);
printf("Result vector is:\n");
for(i=0;i<n;i++)
{
    for(j=0;j<m;j++)
        printf("%d\t",t[i*m+j]);
    printf("\n");
}

getchar();
cudaFree(d_a);
cudaFree(d_t);
return 0;
}

```

Lab Exercises:

- Write a program in CUDA to add two Matrices for the following specifications:
 - Each row of resultant matrix to be computed by one thread.
 - Each column of resultant matrix to be computed by one thread.
 - Each element of resultant matrix to be computed by one thread.
- Write a program in CUDA to multiply two Matrices for the following specifications:
 - Each row of resultant matrix to be computed by one thread.
 - Each column of resultant matrix to be computed by one thread.
 - Each element of resultant matrix to be computed by one thread.

Additional Exercises:

- Write a CUDA program that reads a $M \times N$ matrix A and produces a resultant matrix B of same size as follows: Replace all the even numbered matrix elements with their row sum and odd numbered matrix elements with their column sum.

Example:

	A				B		
I/p: 1	2	3		O/p:	5	6	9
4	5	6			15	7	15

- Write a CUDA program to read a matrix A of size $N \times N$. It replaces the principal diagonal elements with zero. Elements above the principal diagonal by their factorial and elements below the principal diagonal by their sum of digits.

Programs on Matrix using CUDA(contd.)


Objectives:

In this lab, student will be able to

1. Understand how to write kernel code in CUDA to perform operations on matrix.
2. Learn how to represent sparse matrix using CSR format and perform parallel operations using it.

Compressed Sparse Row (CSR) format:

- A sparse matrix is a matrix where the majority of the elements are zero. Sparse matrices are stored in a format that avoids storing zero elements.
- The compressed sparse row (CSR) format is a popular, general-purpose sparse matrix representation. CSR explicitly stores column indices and nonzero values in arrays `col_index` and `data` as shown below.
- A third array of row pointers, `row_ptr`, give the starting location of every row in the compressed storage.



Lab Exercises:

1. Write a program in CUDA to perform parallel Sparse Matrix - Vector multiplication using compressed sparse row (CSR) storage format. Represent the input sparse matrix in CSR format in the host code.
2. Write a program in CUDA to read MXN matrix A and replace 1st row of this matrix by same elements, 2nd row elements by square of each element and 3rd row elements by cube of each element and so on.
3. Write a CUDA program that reads a matrix A of size MXN and produce an output matrix B of same size such that it replaces all the non-border elements(numbers in bold) of A with its equivalent 1's complement and remaining elements same as matrix A.

A				B			
1	2	3	4	1	2	3	4
6	5	8	3	6	10	111	3
2	4	10	1	2	11	101	1
9	1	2	5	9	1	2	5

Additional Exercises:

1. Write a CUDA program which reads an input matrix A of size MXN and produces an output matrix B of size MXN such that, each element of the output matrix is calculated in parallel. Each element, B[i][j], in the output matrix is obtained by adding the elements in ith row and jth column of the input matrix A.

Example:	A			B			
	1	2	3	O/p:	11	13	15
	4	5	6		20	22	24

2. Write a CUDA program that reads a character type matrix A and integer type matrix B of size MXN. It produces an output string *STR* such that, every character of A is repeated *r* times (where *r* is the integer value in matrix B which is having the same index as that of the character taken in A). Write the kernel such that every value of input matrix must be produced required number of times by one thread.

Example:	A				B			
	p	C	a	P	1	2	4	3
	e	X	a	M	2	4	3	2

Output String STR: pCCaaaaPPPeXXXXaaaMM

Programs on Matrix using CUDA and Image Processing **Applications in CUDA**


Objectives:

In this lab, student will be able to,

1. Learn how to deal with matrix using 2D blocks and 2D threads
2. Representing image in raw PGM format.
3. Basic image processing applications and their parallelization.

Image Rotation:

Image rotation is a basic operation in matching, alignment and other image-based algorithms. The input to an image rotation routine is an image, the rotation angle θ , and a point about which a rotation is done.



The coordinates of a point (x_1, y_1) when rotated by an angle y around (x_0, y_0) become (x_2, y_2) , as shown by the following equation:

$$\begin{aligned}x_2 &= \cos(\theta) * (x_1 - x_0) + \sin(\theta) * (y_1 - y_0) \\y_2 &= -\sin(\theta) * (x_1 - x_0) + \cos(\theta) * (y_1 - y_0)\end{aligned}$$

By rotating the image about the origin (0, 0), we get

$$\begin{aligned}x_2 &= \cos(\theta) * x_1 + \sin(\theta) * y_1 \\y_2 &= -\sin(\theta) * x_1 + \cos(\theta) * y_1\end{aligned}$$

JPEG is the widely used image file format. To read the pixels in JPEG, decoding of the image has to be done.

It is easy to read the pixel values of an image, if the pixel values are stored in raw format. Portable Gray Map(PGM) is an image standard.

An example PGM file is as follows.

```
P2
# feep.pgm
24 7
15
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 3 3 3 3 0 0 7 7 7 7 0 0 11 11 11 11 0 0 15 15 15 15 0
0 3 0 0 0 0 0 7 0 0 0 0 0 11 0 0 0 0 0 15 0 0 15 0
0 3 3 3 0 0 0 7 7 7 0 0 0 11 11 11 0 0 0 15 15 15 15 0
0 3 0 0 0 0 0 7 0 0 0 0 0 11 0 0 0 0 0 15 0 0 0 0
0 3 0 0 0 0 0 7 7 7 7 0 0 11 11 11 11 0 0 15 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

The first line P2 indicates plain PGM file.

is a comment line

24 indicates width and 7 indicates height of image.

15 indicates maximum intensity in each pixel.

Rest are intensities in each pixel(total 24*7 pixels).

The above image displays the characters FEEP in increasing order of intensity.

CUDA program for Image rotation

```
#include "cuda_runtime.h"
#include "device_launch_parameters.h"
#include <stdio.h>

#define PI 3.14
```

```

//Fix the DATATYPE to appropriate value
typedef unsigned short DATATYPE;

__global__ void img_rotate(DATATYPE *dest_data, DATATYPE *src_data,int W, int
H, float sinTheta, float cosTheta )
{
//Work-item gets its index within index space
    int ix = threadIdx.x;
    int iy = blockIdx.x;
    //Calculate location of data to move into (ix,iy)
    //Output decomposition as mentioned
    float x0 = W/2.0f;
    float y0=W/2.0f;
    float xOff = ix-x0;
    float yOff = iy-y0;
    int xpos = (int)(xOff*cosTheta + yOff*sinTheta + x0 );
    int ypos = (int)(yOff*cosTheta - xOff*sinTheta + y0 );

    //Bound Checking
    if(((int)xpos>=0) && ((int)xpos< W) &&((int)ypos>=0) && ((int)ypos< H))
    {
        // Read (ix,iy) src_data and store at (xpos,ypos) in
        // dest_data
        // In this case, because we rotating about the origin

        // and there is no translation,
        dest_data[iy*W+ix] = src_data[ypos*W+xpos];
    }
}

struct PGMstructure
{
    int maxVal;
    int width;
    int height;
    DATATYPE *data;
};

```

```

int main()
{
    FILE *imagein,*imageout;
    int row,col;
    int i,j;
    unsigned int ch_int,angle;
    struct PGMstructure *imginfo;
    char fpath[1000],tpath[1000];
    printf("Enter PGM file path:");
    scanf("%s",fpath);
    imagein = fopen(fpath,"r+");
    if(imagein==NULL)
    {
        printf("Error opening first file");
        exit(8);
    }
    while(getc(imagein) != '\n');
    while (getc(imagein) == '#')
    {
        while (getc(imagein) != '\n');
    }
        imginfo->data[row][col] = ch_int;
    }
}

//Allocate memory for structure
imginfo = (struct PGMstructure *)malloc(sizeof(struct PGMstructure));
fseek(imagein, -1, SEEK_CUR);
fscanf(imagein,"%d", &imginfo->height);
fscanf(imagein,"%d", &imginfo->width);
fscanf(imagein,"%d", &imginfo->maxVal);
printf("\n width = %d\n",imginfo->width);
printf("\n height = %d\n",imginfo->height);

printf("\n maxVal = %d\n",imginfo->maxVal);
imginfo->data=(DATATYPE *)calloc(imginfo->width*imginfo-
>height,sizeof(DATATYPE *));
for (row=0;row<imginfo->height; row++)
{
    //imginfo->data[row*imginfo->width]=(DATATYPE *)calloc(imginfo-
>width,sizeof(DATATYPE));

```

```

        for (col=0; col< imginfo->width; col++)
        {
            fscanf(imagein,"%u", &ch_int);
            imginfo->data[row*imginfo->width+col] = ch_int;
        }
    }
    printf("Enter the angle to rotate");
    scanf("%d",&angle);
    float sinTheta=sinf(angle*PI/180.0f);
    float cosTheta=cosf(angle*PI/180.0f);

    DATATYPE *d_src, *d_dst, *dst;
    int size = imginfo->width*imginfo->height* sizeof(DATATYPE);

    // Alloc space for device copies of d_src,d_dst
    cudaMalloc((void **)&d_src, size);
    cudaMalloc((void **)&d_dst, size);

    cudaMemcpy(d_src, imginfo->data, size, cudaMemcpyHostToDevice);

    // Launch kernel on GPU
    img_rotate<<<imginfo->width,imginfo->height>>>(d_dst,d_src,imginfo->
width,imginfo->height,          sinTheta,cosTheta);

    dst=(DATATYPE *)malloc(imginfo->width*imginfo->
height*sizeof(DATATYPE));

    // Copy result back to host
    cudaMemcpy(dst, d_dst, size, cudaMemcpyDeviceToHost);
    printf("Enter path of output file:");

    scanf("%s",tpath);
    imageout=fopen(tpath,"w+");

    //Write the PGM file header
    fputs("P2\n",imageout);

    //Write the height and width
    fprintf(imageout,"%d\n",imginfo->height);

```

```

fprintf(imageout,"%d\n",imginfo->width);
fprintf(imageout,"%d\n",imginfo->maxVal);

for ( i = 0 ; i < imginfo->height ; i++ )
{
    for ( j = 0 ; j < imginfo->width ; j++ )
        fprintf( imageout,"%d\n" , dst[i*imginfo->width+j] );
}
fclose(imageout);


// Cleanup
free(imginfo->data);
free(dst);
cudaFree(d_src);
cudaFree(d_dst);
return 0;
}

```

Image Convolution:

In image processing, convolution is a commonly used algorithm that modifies the value of each pixel in an image by using information from neighboring pixels. A convolution kernel, or filter, describes how each pixel will be influenced by its neighbors. For example, a blurring kernel will take the weighted average of neighboring pixels so that large differences between pixel values are reduced. By using the same source image and changing only the filter, effects such as sharpening, blurring, edge enhancing, and embossing can be produced.

A convolution kernel works by iterating over each pixel in the source image. For each source pixel, the filter is centered over the pixel and the values of the filter multiply the pixel values that they overlay. A sum of the products is then taken to produce a new pixel value. Following figure provides a visual for this algorithm.



Solved Exercise:

Write a program in CUDA to perform matrix addition using 2D Grid and 2D Block

```
//Matrix addition of 4x4 matrix
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <cuda_runtime.h>
#define BLOCK_WIDTH 2
#define TILE_WIDTH 2
#define WIDTH 4

__device__ int getTid() {
    int blockSkip = (blockIdx.y * gridDim.x * blockDim.x * blockDim.y);
    int rowSkip = (threadIdx.y * gridDim.x * blockDim.x);
    int rowDisplacement = (blockIdx.x * blockDim.x) + threadIdx.x;
    int tid = blockSkip + rowSkip + rowDisplacement;
    return tid;
}

__global__ void MatAddElementThread(int *a, int *b, int *d) {
    int tid = getTid();
    d[tid] = a[tid] + b[tid];
}

int main() {
    int *matA, *matB, *matSum;
    int *da, *db, *dc;

    printf("\n== Enter elements of Matrix A (4x4) ==\n");

    matA = (int*)malloc(sizeof(int) * WIDTH * WIDTH);
    for(int i = 0; i < WIDTH * WIDTH; i++)
    {
        scanf("%d", &matA[i]);
    }

    printf("\n== Enter elements of Matrix B (4x4) ==\n");
    matB = (int*)malloc(sizeof(int) * WIDTH * WIDTH);
```

```

for(int i = 0; i < WIDTH * WIDTH; i++)
{
    scanf("%d", &matB[i]);
}
matSum = (int*)malloc(sizeof(int) * WIDTH * WIDTH);

cudaMalloc((void **) &da, sizeof(int) * WIDTH * WIDTH);
cudaMalloc((void **) &db, sizeof(int) * WIDTH * WIDTH);
cudaMalloc((void **) &dc, sizeof(int) * WIDTH * WIDTH);

cudaMemcpy(da, matA, sizeof(int) * WIDTH * WIDTH, cudaMemcpyHostToDevice);
cudaMemcpy(db, matB, sizeof(int) * WIDTH * WIDTH, cudaMemcpyHostToDevice);
int NumBlocks = WIDTH / BLOCK_WIDTH;
dim3 grid_conf (NumBlocks, NumBlocks);
dim3 block_conf (BLOCK_WIDTH, BLOCK_WIDTH);

    MatAddElementThread<<<grid_conf, block_conf>>>(da, db, dc);

cudaMemcpy(matSum,dc,sizeof(int)* WIDTH *WIDTH,cudaMemcpyDeviceToHost);
printf("\n-=Result of Addition=-\n");
printf(".....\n");
for (int i = 0; i < m; i++ ) {
    for (int j = 0; j < n; j++) {
        printf("%6d ", matSum[i * n + j]);
    }
    printf("\n");
}
cudaFree(da);
cudaFree(db);
cudaFree(dc);
free(matA);
free(matB);
free(matSum);
return 0;
}

```


Lab Exercises :

1. Write a program in CUDA to perform matrix multiplication using 2D Grid and 2D Block.
2. Write a CUDA program for rotating an image after reading the image. The image is in pgm format.

Additional Excersices:

1. Write a CUDA program for reversing the gray-level of an image. Assume the image is in pgm format.
2. Write a CUDA program for averaging filter on an image in pgm format using convolution. The 3X3 mask has center value 8 and all other masks have -1.
3. Write a program in CUDA to perform 2D convolution which takes 2D input array and 2D mask array to produce 2D output array.

Programs using different CUDA Device memory types and synchronization

Synchronization: CUDA allows threads in the same block to coordinate their activities using a barrier synchronization function, `__syncthreads()`. When a kernel function calls `__syncthreads()`, the thread that executes the function call will be held at the calling location until every thread in the block reaches the location. This ensures that all threads in a block have completed a phase of their execution of the kernel before any moves on to the next phase.

Shared Variables: Accessing shared memory is extremely fast and highly parallel. If a variable declaration is preceded by the keyword `__shared__`, it declares a shared variable in CUDA. Such declarations typically reside within a kernel function or a device function. The scope of a shared variable is within a thread block means all threads in a block see the same version of a shared variable. The lifetime of a shared variable is within the duration of the kernel. Shared variables are an efficient means for threads within a block to collaborate with each other. Shared memory is fast but it is small. A common strategy is partition the data into subsets called **tiles** so that each tile fits into the shared memory.

Constant Variables: If a variable declaration is preceded by the keyword `__constant__`, it declares a constant variable in CUDA. Declaration of constant variables must be outside any function body. The scope of a constant variable is all grids, meaning that all threads in all grids see the same version of a constant variable. The lifetime of a constant variable is the entire application execution. Constant variables are stored in the global memory but are cached for efficient access. With appropriate access patterns, accessing constant memory is extremely fast and parallel. Currently, the total size of constant variables in an application is limited at 65,536 bytes. One may need to break up the input data volume to fit within this limitation.

Device Variables: A variable whose declaration is preceded only by the keyword `__device__` is a global variable and will be placed in global memory. Accesses to a global variable are slow. However, global variables are visible to all threads of all kernels. Their contents also persist through the entire execution. Thus, global variables can be used as a means for threads to collaborate across blocks. Global variables are often used to pass information from one kernel invocation to another kernel invocation.

Solved Exercise:

Write a program in CUDA to perform tiled matrix multiplication using 2D Grid and 2D Block

```
//Matrix multiplication of 4x4 matrix
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <cuda_runtime.h>
#define BLOCK_WIDTH 2
#define TILE_WIDTH 2
#define WIDTH 4

__global__ void MatMulElementThreadShared(int *a, int *b, int *c) {
    __shared__ int MDs[TILE_WIDTH][TILE_WIDTH];
    __shared__ int NDs[TILE_WIDTH][TILE_WIDTH];
    int m;
    int bx=blockIdx.x; int by=blockIdx.y;
    int tx=threadIdx.x; int ty=threadIdx.y;

    int Row=by*TILE_WIDTH + ty;
    int Col= bx*TILE_WIDTH + tx;

    int Pvalue=0;
    for(m=0; m<WIDTH/TILE_WIDTH; m++)
    {
        MDs[ty][tx]=a[Row*WIDTH+m*TILE_WIDTH+tx];
        NDs[ty][tx]=b[(m*TILE_WIDTH+ty)*WIDTH+Col];

        __syncthreads();

        for (int k = 0; k < TILE_WIDTH; k++)
        {
            Pvalue += MDs[ty][k]*NDs[k][tx];
        }
        __syncthreads();
    }
    c[Row*WIDTH+Col] = Pvalue;
}
```

```

int main() {
    int *matA, *matB, *matProd;
    int *da, *db, *dc;

    printf("\n== Enter elements of Matrix A (4x4) ==\n");

    matA = (int*)malloc(sizeof(int) * WIDTH * WIDTH);
    for(int i = 0; i < WIDTH * WIDTH; i++)
    {
        scanf("%d", &matA[i]);
    }

    printf("\n== Enter elements of Matrix B (4x4) ==\n");
    matB = (int*)malloc(sizeof(int) * WIDTH * WIDTH);
    for(int i = 0; i < WIDTH * WIDTH; i++)
    {
        scanf("%d", &matB[i]);
    }
    matProd = (int*)malloc(sizeof(int) * WIDTH * WIDTH);

    cudaMalloc((void **) &da, sizeof(int) * WIDTH * WIDTH);
    cudaMalloc((void **) &db, sizeof(int) * WIDTH * WIDTH);
    cudaMalloc((void **) &dc, sizeof(int) * WIDTH * WIDTH);

    cudaMemcpy(da, matA, sizeof(int) * WIDTH * WIDTH, cudaMemcpyHostToDevice);
    cudaMemcpy(db, matB, sizeof(int) * WIDTH * WIDTH, cudaMemcpyHostToDevice);
    int NumBlocks = WIDTH / BLOCK_WIDTH;
    dim3 grid_conf (NumBlocks, NumBlocks);
    dim3 block_conf (BLOCK_WIDTH, BLOCK_WIDTH);

    MatMulElementThreadShared<<<grid_conf, block_conf>>>(da, db, dc);

    cudaMemcpy(matProd, dc, sizeof(int) * WIDTH * WIDTH, cudaMemcpyDeviceToHost);
    printf("\n==Result of Addition==\n");
    printf(".....\n");
    for (int i = 0; i < m; i++ ) {
        for (int j = 0; j < n; j++) {
            printf("%6d ", matProd[i * n + j]);
        }
    }
}

```

```
        printf("\n");
    }
    cudaFree(da);
    cudaFree(db);
    cudaFree(dc);
    free(matA);
    free(matB);
    free(matProd);
    return 0;
}
```

Lab Exercises:

1. Write a program in CUDA to improve the performance of 1D parallel convolution using constant Memory.
2. Write a program in CUDA to perform tiled 1D convolution operation on the input array N of size *width* using the mask array, M of size *mask_width*, to produce the resultant array P of size *width*.
3. Write a program in CUDA to perform inclusive scan algorithm.

Additional Exercise:

1. Write a program in CUDA which displays a shopping mall item menu with its price. The N number of friends are allowed to purchase as many items they want. Calculate the total purchase done by N friends.

REFERENCES

1. Michael J. Quinn, “*Parallel Programming in C with MPI and OpenMP*”, McGraw Hill Edition, 2003.
2. Benedict R. Gaster, Lee Howes, David R, Perhaad Mistry, Dana Schaa, “*Heterogeneous Computing with OpenCL*”, Elsevier Inc., 1st Edition, 2012.
3. D. Kirk and W. Hwu , “*Programming Massively Parallel Processors –A Hands-on approach*”, Elsevier Inc., 2nd Edition, 2013.
4. Gonzalez, Rafael C., and Richard E. Woods. "Digital image processing [M]." *Publishing house of electronics industry*141.7 (2002).